

Course 5: Reinforcement Learning



IMT Atlantique
Bretagne-Pays de la Loire
École Mines-Télécom

Last session

- 1 Combinatorial game theory
- 2 Definition of a game
- 3 Proof of determined games

Today's session

- 1 Reinforcement Learning
- 2 Value and Policy Functions
- 3 Q-Learning

Note: reinforcement learning and combinatorial game theory share a common mathematical framework. But to ease access to online resources, we will adopt a new vocabulary.

Outline of the course

1 Definitions of Reinforcement Learning (RL)

- Fundamentals
- Example: PyRat
- Policy and values

2 Q-learning

- Q-learning definitions
- Example
- Approximate Q-learning
- Exploration/Exploitation

1 Definitions of Reinforcement Learning (RL)

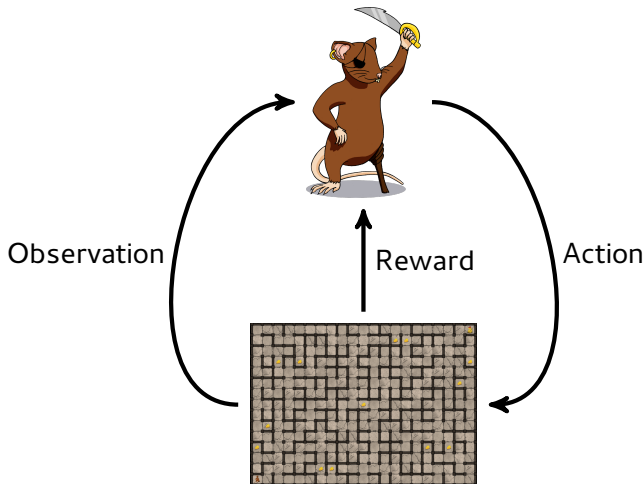
- Fundamentals
- Example: PyRat
- Policy and values

2 Q-learning

- Q-learning definitions
- Example
- Approximate Q-learning
- Exploration/Exploitation

Agent and environment

Our objective is to train an **agent** to maximize its **reward** through **actions** that affect an **environment**.



Reward hypothesis

- *All goals can be described by the maximization of expected cumulated reward over time.*

Specificities of reinforcement learning

- No supervision, only a reward signal,
- Delayed feedback, the reward can come (much) later,
- Importance of the temporal dimension,
- Agent's actions affect the subsequent data it receives.

Reward hypothesis


- *All goals can be described by the maximization of expected cumulated reward over time.*

Specificities of reinforcement learning

- No supervision, only a reward signal,
- Delayed feedback, the reward can come (much) later,
- Importance of the temporal dimension,
- Agent's actions affect the subsequent data it receives.

Agent and environment

Agent state representation $s_t^\alpha \in \mathcal{S}^\alpha$



Environment state representation $s_t^e \in \mathcal{S}^e$

Definitions

■ The agent α ...

- 1 analyzes previous actions, states, rewards and observations,
- 2 computes action a_t ,
- 3 obtains reward r_t ,
- 4 obtains an observation o_{t+1} ,
- 5 deduce a new state s_{t+1}^α .

■ The environment...

- 1 receives action a_t ,
- 2 produces reward r_t ,
- 3 deduce a new state s_{t+1}^e ,
- 4 produces o_{t+1} .

■ Observability:

- **Perfect:** $s_t^\alpha = s_t^e = o_t$
- **Imperfect:** no access to full environment state:
 - The agent indirectly observes the environment through o_t ,
 - s_t^α is estimated by the agent and may differ from s_t^e .

Definitions

■ The agent α ...

- 1 analyzes previous actions, states, rewards and observations,
- 2 computes action a_t ,
- 3 obtains reward r_t ,
- 4 obtains an observation o_{t+1} ,
- 5 deduce a new state s_{t+1}^α .

■ The environment...

- 1 receives action a_t ,
- 2 produces reward r_t ,
- 3 deduce a new state s_{t+1}^e ,
- 4 produces o_{t+1} .

■ Observability:

- **Perfect:** $s_t^\alpha = s_t^e = o_t$
- **Imperfect:** no access to full environment state:
 - The agent indirectly observes the environment through o_t ,
 - s_t^α is estimated by the agent and may differ from s_t^e .

Definitions

■ The agent α ...

- 1 analyzes previous actions, states, rewards and observations,
- 2 computes action a_t ,
- 3 obtains reward r_t ,
- 4 obtains an observation o_{t+1} ,
- 5 deduce a new state s_{t+1}^α .

■ The environment...

- 1 receives action a_t ,
- 2 produces reward r_t ,
- 3 deduce a new state s_{t+1}^e ,
- 4 produces o_{t+1} .

■ Observability:

- **Perfect:** $s_t^\alpha = s_t^e = o_t$
- **Imperfect:** no access to full environment state:
 - The agent indirectly observes the environment through o_t ,
 - s_t^α is estimated by the agent and may differ from s_t^e .

Definitions

■ The agent α ...

- 1 analyzes previous actions, states, rewards and observations,
- 2 computes action a_t ,
- 3 obtains reward r_t ,
- 4 obtains an observation o_{t+1} ,
- 5 deduce a new state s_{t+1}^α .

■ The environment...

- 1 receives action a_t ,
- 2 produces reward r_t ,
- 3 deduce a new state s_{t+1}^e ,
- 4 produces o_{t+1} .

■ Observability:

- **Perfect:** $s_t^\alpha = s_t^e = o_t$
- **Imperfect:** no access to full environment state:
 - The agent indirectly observes the environment through o_t ,
 - s_t^α is estimated by the agent and may differ from s_t^e .

Example: PyRat

Definitions

- The agent is either the Rat or the Python,
- The opponent becomes part of the environment,
 - Note that the game can be with perfect observability if the opponent strategy is known,
- Seen this way, the game becomes sequential.

RL-based PyRat versus supervised approach

- Reward signal: number of picked up pieces of cheese,
- Delayed feedback: several moves required to reach a reward,
- Character's moves affect subsequent data it receives,
- Importance of the temporal dimension.

Example: PyRat

Definitions

- The agent is either the Rat or the Python,
- The opponent becomes part of the environment,
 - Note that the game can be with perfect observability if the opponent strategy is known,
- Seen this way, the game becomes sequential.

RL-based PyRat versus supervised approach

- Reward signal: number of picked up pieces of cheese,
- Delayed feedback: several moves required to reach a reward,
- Character's moves affect subsequent data it receives,
- Importance of the temporal dimension.

Observability examples

- Perfect: $s_t^{rat} = s_t^o = o_t$
 - a_t : Last move of the rat,
 - o_t : The entire maze with all cheese locations and python position,
 - r_t : Binary variable which is 1 if the rat just got a piece of cheese.
- Imperfect:
 - a_t : Last move of the rat,
 - o_t : Neighboring cells of the rat,
 - r_t : Binary variable which is 1 if the rat just got a piece of cheese.

To represent the strategy of the rat, we use a **policy function**.

Observability examples

- Perfect: $s_t^{rat} = s_t^o = o_t$
 - a_t : Last move of the rat,
 - o_t : The entire maze with all cheese locations and python position,
 - r_t : Binary variable which is 1 if the rat just got a piece of cheese.
- Imperfect:
 - a_t : Last move of the rat,
 - o_t : Neighboring cells of the rat,
 - r_t : Binary variable which is 1 if the rat just got a piece of cheese.

To represent the strategy of the rat, we use a **policy function**.

Observability examples

- Perfect: $s_t^{rat} = s_t^o = o_t$
 - a_t : Last move of the rat,
 - o_t : The entire maze with all cheese locations and python position,
 - r_t : Binary variable which is 1 if the rat just got a piece of cheese.
- Imperfect:
 - a_t : Last move of the rat,
 - o_t : Neighboring cells of the rat,
 - r_t : Binary variable which is 1 if the rat just got a piece of cheese.

To represent the strategy of the rat, we use a [policy function](#).

Policy Function

Definition

The policy function of an agent α is:

$$\pi : \begin{cases} S^\alpha & \rightarrow A \\ s_t^\alpha & \mapsto a_t^\alpha \end{cases}$$

- π can be deterministic or stochastic.

Playout

The playout $(s_t^{\alpha, \pi})_{t \in \mathbb{N}}$ associated with a policy π and initial state s_0 , is defined by considering agent α takes his/her actions using π .

Policy Function

Definition

The policy function of an agent α is:

$$\pi : \begin{cases} S^\alpha & \rightarrow A \\ s_t^\alpha & \mapsto a_t^\alpha \end{cases}$$

- π can be deterministic or stochastic.

Playout

The playout $(s_t^{\alpha, \pi})_{t \in \mathbb{N}}$ associated with a policy π and initial state s_0 , is defined by considering agent α takes his/her actions using π .

Value Function

Definition

Fix $\gamma \in [0, 1[$, the value function v^π is defined as:

$$v^\pi : \begin{cases} \mathcal{S}^\alpha & \rightarrow \mathbb{R} \\ s_{t_0}^{\alpha, \pi} & \mapsto \sum_{t=t_0}^{+\infty} \gamma^{t-t_0} r_t \end{cases}$$

- The value of a policy function is thus an expectation of cumulative future rewards, weakened by the geometrical coefficient γ to avoid divergence,
- The best possible policy $\pi^*(s)$ is defined as:

$$\forall s \in \mathcal{S}^\alpha, \forall \pi, V^{\pi^*}(s) \geq V^\pi(s).$$

Value Function

Definition

Fix $\gamma \in [0, 1[$, the value function v^π is defined as:

$$v^\pi : \begin{cases} \mathcal{S}^\alpha & \rightarrow \mathbb{R} \\ s_{t_0}^{\alpha, \pi} & \mapsto \sum_{t=t_0}^{+\infty} \gamma^{t-t_0} r_t \end{cases}$$

- The value of a policy function is thus an expectation of cumulative future rewards, weakened by the geometrical coefficient γ to avoid divergence,
- The best possible policy $\pi^*(s)$ is defined as:

$$\forall s \in \mathcal{S}^\alpha, \forall \pi, V^{\pi^*}(s) \geq V^\pi(s).$$

1 Definitions of Reinforcement Learning (RL)

- Fundamentals
- Example: PyRat
- Policy and values

2 Q-learning

- Q-learning definitions
- Example
- Approximate Q-learning
- Exploration/Exploitation

Q-learning

Definition

In Q-learning, we aim to find V^{π^*} as a solution to the recursive system of equations (Bellman equation):

$$\forall s \in \mathcal{S}^\alpha, \forall a \in A, Q(s, a) = r_{s,a} + \gamma \max_{a'} Q(s(a), a'),$$

where $r_{s,a}$ is the reward agent α performs action a in state s and $s(a)$ is the state observed by agent α after performing action a .

Pros and cons

Pros:

- Can be learned even if the agent is not following any specific π ,
- Self training is possible,

Cons:

- Scalability issues when \mathcal{S}^α is large.

Q-learning

Definition

In Q-learning, we aim to find V^{π^*} as a solution to the recursive system of equations (Bellman equation):

$$\forall s \in \mathcal{S}^\alpha, \forall a \in A, Q(s, a) = r_{s,a} + \gamma \max_{a'} Q(s(a), a'),$$

where $r_{s,a}$ is the reward agent α performs action a in state s and $s(a)$ is the state observed by agent α after performing action a .

Pros and cons

Pros:

- Can be learned even if the agent is not following any specific π ,
- Self training is possible,

Cons:

- Scalability issues when \mathcal{S}^α is large.

Q-learning

Definition

In Q-learning, we aim to find V^{π^*} as a solution to the recursive system of equations (Bellman equation):

$$\forall s \in \mathcal{S}^\alpha, \forall a \in A, Q(s, a) = r_{s,a} + \gamma \max_{a'} Q(s(a), a'),$$

where $r_{s,a}$ is the reward agent α performs action a in state s and $s(a)$ is the state observed by agent α after performing action a .

Pros and cons

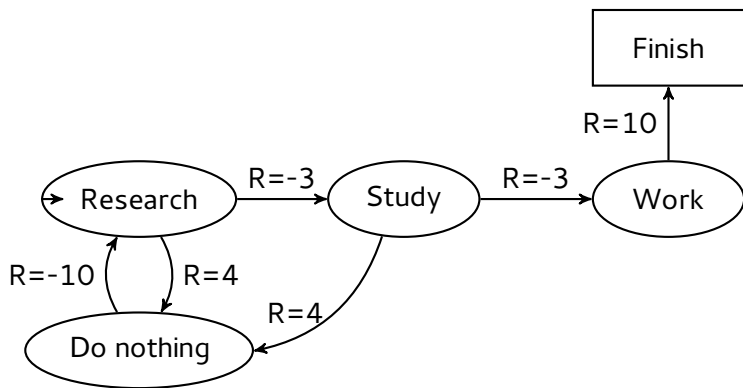
Pros:

- Can be learned even if the agent is not following any specific π ,
- Self training is possible,

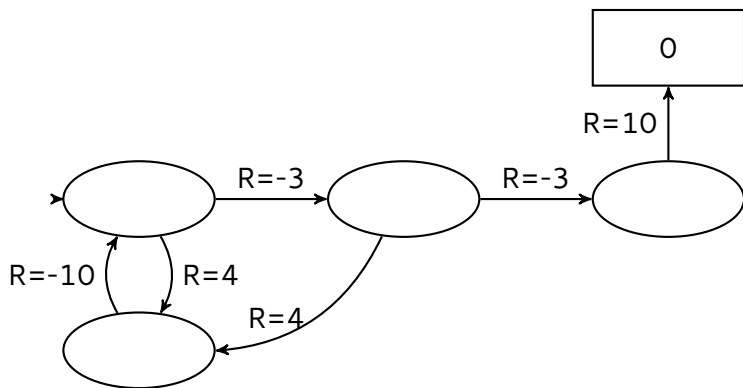
Cons:

- Scalability issues when \mathcal{S}^α is large.

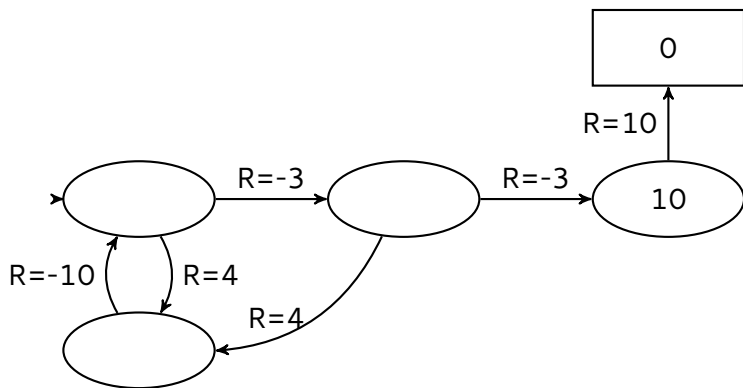
Q-learning example



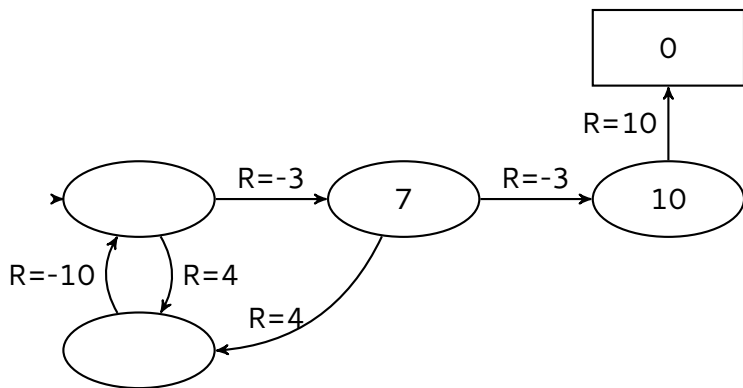
Q-learning example



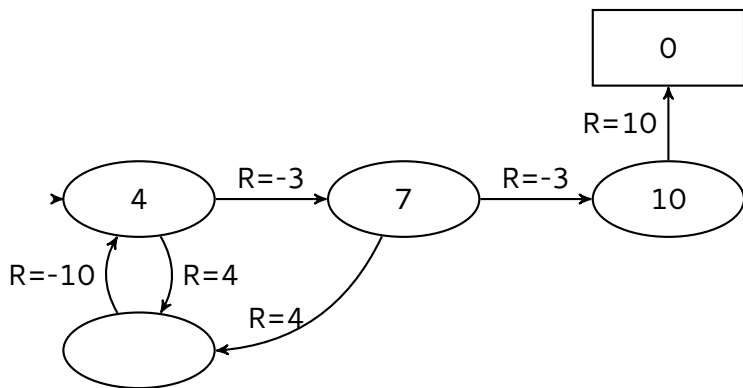
Q-learning example



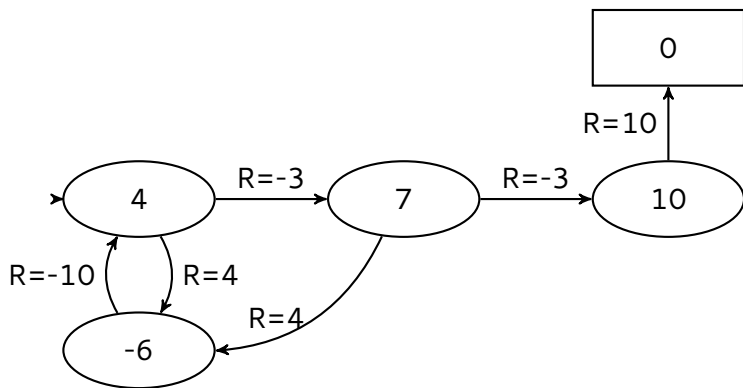
Q-learning example




Q-learning example



Q-learning example



Q-learning example



Approximated Q-learning

Definition

- Train a model to approximate Q ,
 - Input is a state s and output is made of values of $Q(s, \cdot)$,
 - Representation learning can be used to compress S^α .

Problems

- Almost always needs a simulator for the game,
- Game duration can be bottleneck for training,
- Catastrophic forgetting and adversary specialization,
 - These effects can be alleviated by training using experience replay.

Experience replay

- Instead of using only the last decision to train, sample at random from the m previous decisions,
- Decisions taken before should remain considered now.

Approximated Q-learning

Definition

- Train a model to approximate Q ,
 - Input is a state s and output is made of values of $Q(s, \cdot)$,
 - Representation learning can be used to compress S^α .

Problems

- Almost always needs a simulator for the game,
- Game duration can be bottleneck for training,
- Catastrophic forgetting and adversary specialization,
 - These effects can be alleviated by training using experience replay.

Experience replay

- Instead of using only the last decision to train, sample at random from the m previous decisions,
- Decisions taken before should remain considered now.

Approximated Q-learning

Definition

- Train a model to approximate Q ,
 - Input is a state s and output is made of values of $Q(s, \cdot)$,
 - Representation learning can be used to compress S^α .

Problems

- Almost always needs a simulator for the game,
- Game duration can be bottleneck for training,
- Catastrophic forgetting and adversary specialization,
 - These effects can be alleviated by training using experience replay.

Experience replay

- Instead of using only the last decision to train, sample at random from the m previous decisions,
- Decisions taken before should remain considered now.

Exploration/Exploitation

Dilemma

- Repeat with existing strategy (Exploitation)...
- ... or try a new strategy (Exploration)?

Example

- Always eating in restaurants that you know is exploitation,
- While that is a good heuristic, you have no way of knowing if you have the maximum reward possible,
- So exploring new restaurants from time to time may be needed to find the maximum reward.

Exploration/Exploitation

Dilemma

- Repeat with existing strategy (Exploitation)...
- ... or try a new strategy (Exploration)?

Example

- Always eating in restaurants that you know is exploitation,
- While that is a good heuristic, you have no way of knowing if you have the maximum reward possible,
- So exploring new restaurants from time to time may be needed to find the maximum reward.

TP4 - PyRat with reinforcement learning

- Approximate Q-learning algorithm using experience replay and linear regression to beat the greedy algorithm,
- Approximation method (linear regression) and experience replay routine are given,
- Assemble all the primitives to perform Reinforcement Learning.

Challenge

You can continue working in the challenge after finishing TP4. You can now integrate reinforcement learning in your solution.